

Discrete Simulations Project

SYSC 4005

By. Emmitt Luning, Chase Scott

Github: <https://github.com/Em-kale/discrete-simulations-project>

Problem Formulation

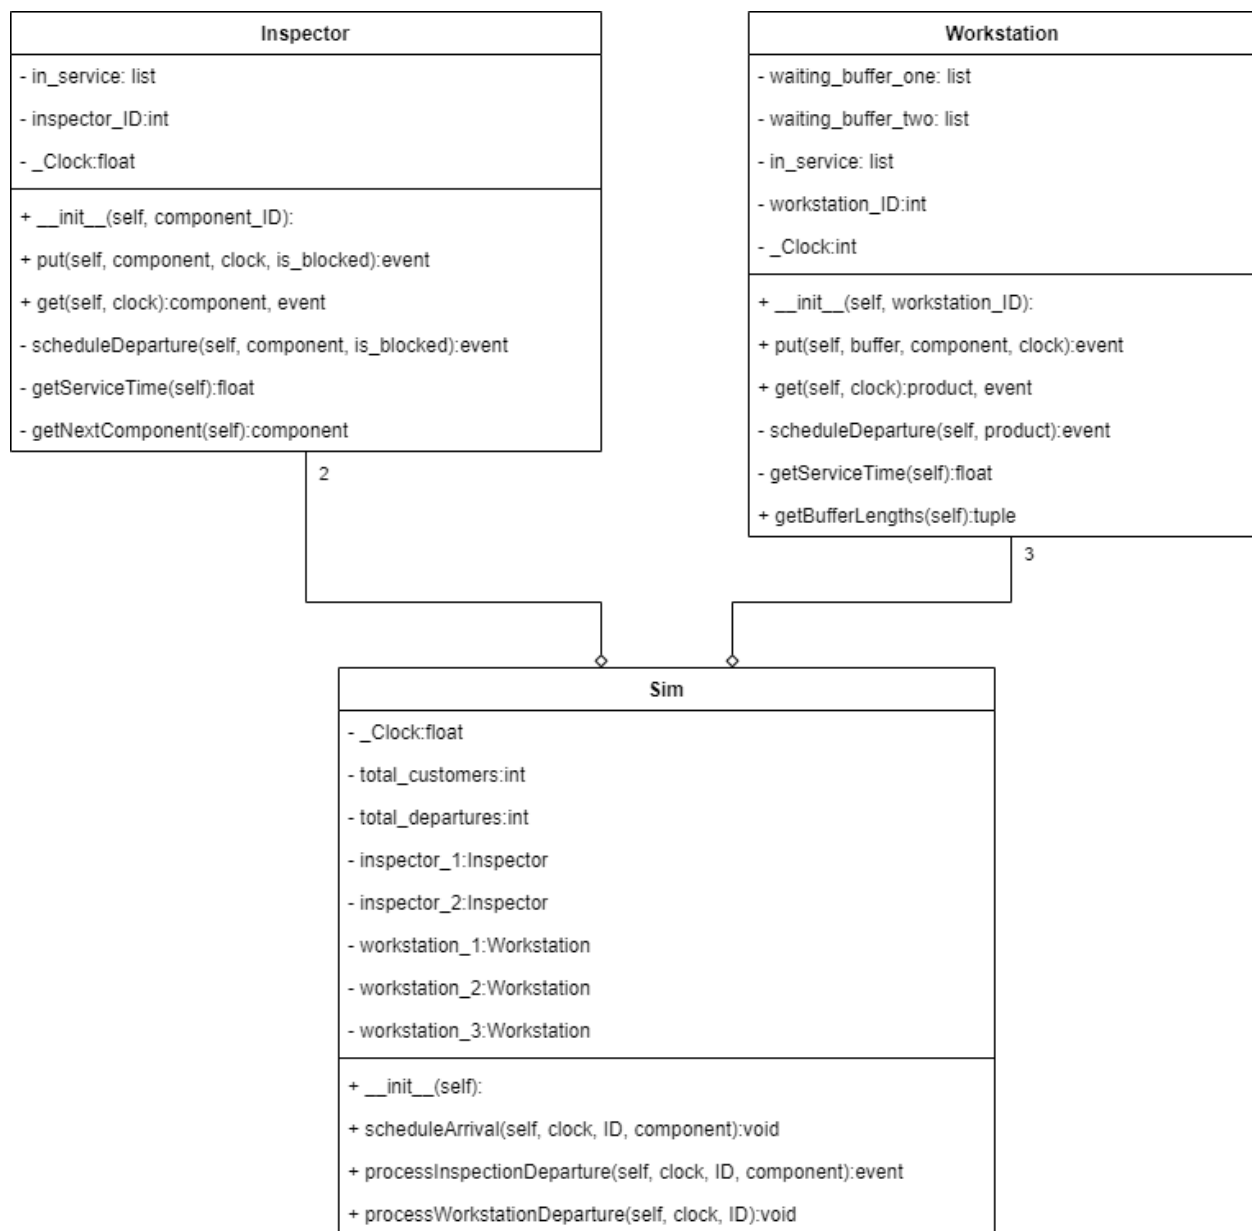
Manufacturing Facility Entities

Within the manufacturing facility there exist two types of stations, inspectors and workstations, along with two types of work products, components and products. The inspectors inspect components, while the workstations assemble products from components that have already been inspected. The components come in three variations denoted as C1, C2, and C3, while the products come in three variations denoted as P1, P2, and P3. Products are constructed of components following the scheme:

- P1: 1 C1 component
- P2: 1 C1 component, and 1 C2 component
- P3: 1 C1 component, and 1 C3 component

Each workstation contains a buffer of size 2 for each type of product that it receives. For example, W3, which assembles products of type P3, has a buffer for C1 components, and a buffer for C3 components, each with a maximum length of two.

The purpose of this simulation is to determine the performance of the manufacturing facility, as well as to explore alternatives for delivering components to workstations from inspector 1.



Setting of Objectives

Simulation Suitability

To determine whether a simulation is adequate to solve the specified problem, consider the metrics of complexity, historical data, verification/validation efficacy, cost, difficulty of direct experimentation, and time.

The specified problem is not too complex to model, which indicates that a simulation would be possible to implement. Additionally, historical data has been provided for which analysis can be

done to determine distributions for input domain modeling, and values can therefore be generated relatively easily. Furthermore, as the simulation would be of a manufacturing facility, there is no doubt that it would be cheaper, and more time effective to simulate the process, than to halt manufacturing and perform direct experimentation, where the time of manufacturing would come into play, as well as the costs associated with manufacturing the goods.

Metrics for Evaluation

Several metrics will be used to evaluate the performance of the system. The amount of products delivered by the system in a given unit of time (throughput) will be calculated along with the average capacity of the workstation buffers, the amount of time the workstations are busy, and the blocked time of the inspectors.

The goal of the manufacturing facility should be to increase the throughput, along with the average buffer capacity, and workstation busy time, while minimizing the buffer time. These goals are how the performance of the system will be evaluated based on the data collected from the simulation.

Resources, and Schedule

This simulation project will have three developers dedicated to the formulation of the simulation, and the subsequent analysis of the data. The timeframe for completion is two months, separated into several milestones as follow:

Input modeling and Generation

The purpose of this milestone will be to find a suitable distribution to use for simulating service times for the inspectors, and workstations, as well as implementing that solution in the simulation.

Model Verification and Validation

The purpose of this milestone will be to validate the model using a variety of verification techniques, and to run the simulation to compute the values of interest, along with their confidence intervals and justification for the initialization length.

Alternative Operating Policies

This milestone will use the gathered information to propose alternative policies to increase the values discussed that should be maximized, and minimize those which should be minimized, as discussed in the Metrics for Evaluation section.

Model Conceptualization

Stakeholders

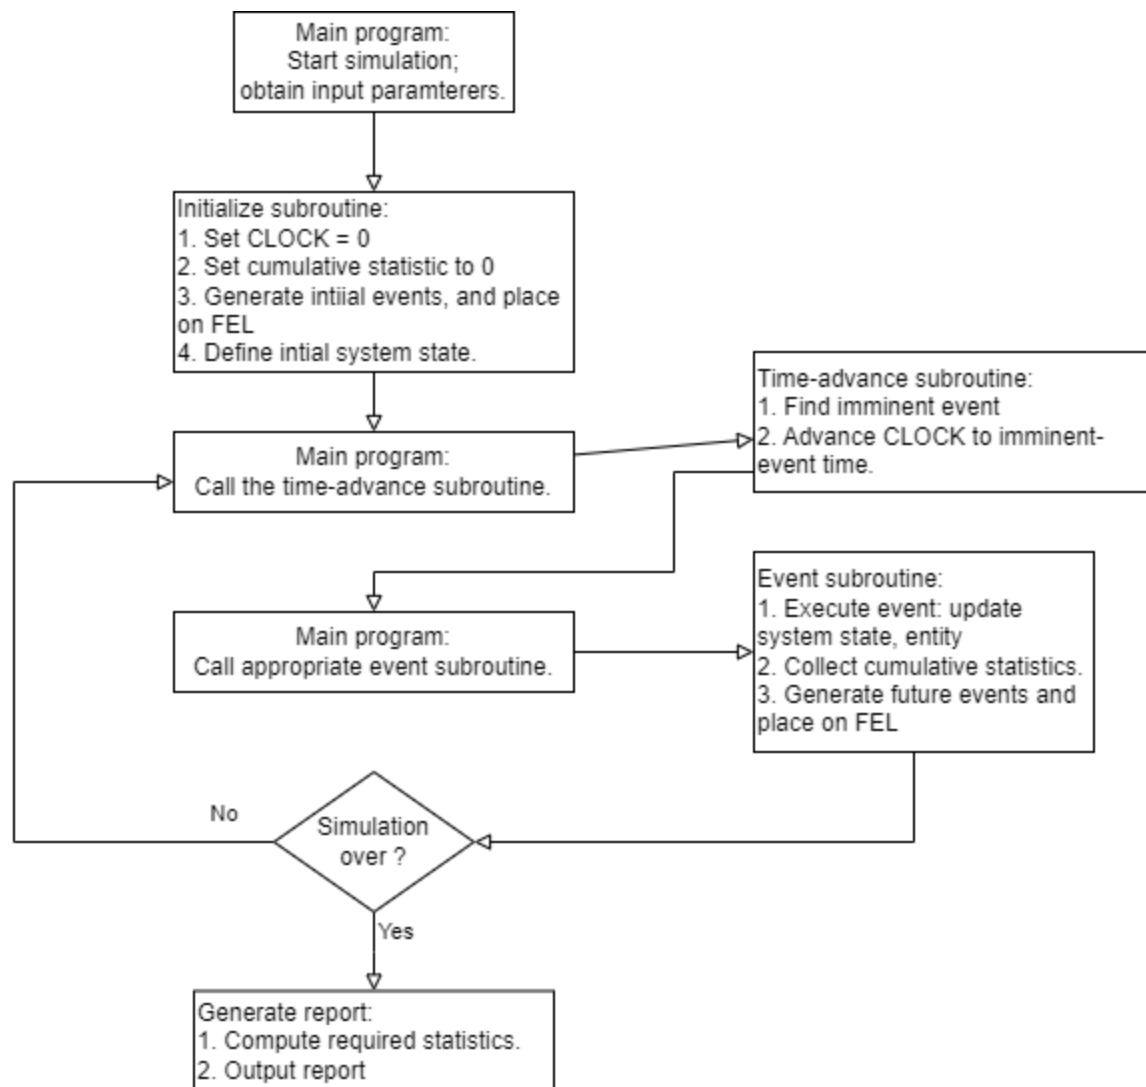
The identified stakeholders for this simulation project are the operators of the manufacturing facility being analyzed. In order for the simulation to be successful it must address their needs by supplying the performance information requested.

Entities and Interactions

The simulation will consist of three major components, the simulation controller itself, the inspectors, and the workstations. These will be represented in the implementation as three distinct classes.

Main Loop

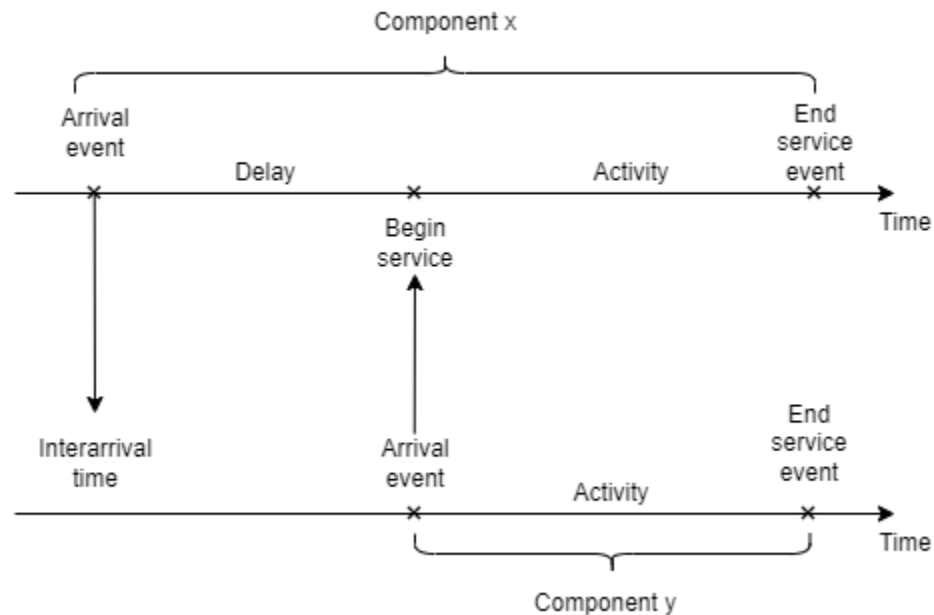
The driving force of the simulation will be the primary loop, which removes an event from the future event list (FEL), executes some subroutine based on the type of event it is, which may or may not place new events in the FEL, and then repeats until a set number of departures from the system have been met. Here, the logic for a basic event loop utilizing an FEL can be seen, this is the basis of the system we are going to implement.



Simulation Class

The simulation class represents the simulation as a whole, and will instantiate two inspectors, and three workstations from their respective classes. It then includes two functions, schedule arrival, and schedule departure, which will handle the two primary types of events. Depending on the type of event, and the current state of the system these functions will make calls to update the inspectors, and workstations to reflect the event that has occurred. This class also contains the logic which decides where to send the components to based on the status of the workstation buffers.

Here shows the interaction for a workstation, a component arrives at a certain time, then is delayed until the other component arrives. Then the workstation can assemble the product from these components:



Inspector Class

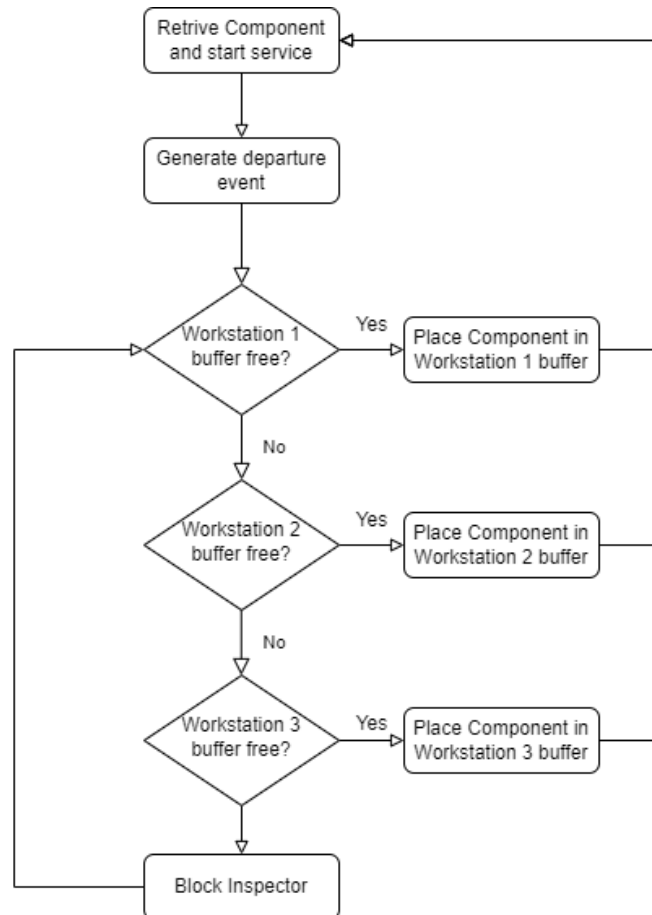
The inspector class represents the inspectors in the simulation. They include two primary functions that are accessed by the simulation class. The primary functions are put, and get, which are used to add new components to the inspector, and try to remove components from the inspector respectively. The put function is only called when initializing the system, as all consequent movement of components from waiting to in service will be done upon the previous component being removed, and this action is performed by the get method. The put method also includes a flag to indicate if the component is currently in service, or if the inspector is currently blocked, which will result in a new departure event being generated for a short period of time in the future.

```
def put(self, component, clock, is_blocked):  
    """ update clock """  
    self._Clock = clock  
    """ start service """  
    self.in_service.append(component)  
    depart = self.scheduleDeparture(component, is_blocked)  
    return depart
```

Additionally, the inspector class includes a method for generating departure events which are returned to the simulation after a new component is added to the inspector and include the necessary information for the simulation to add these to the FEL.

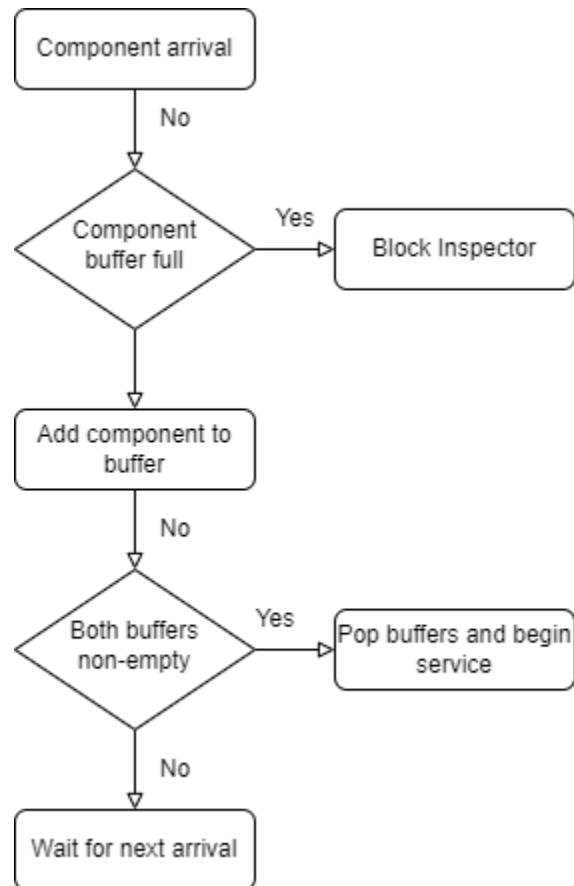
```
def get(self, clock):  
    """ get component from the service queue"""  
    component = self.in_service.pop(0)  
    """ update clock"""  
    self._Clock = clock  
    """ get next component to be serviced"""  
    next_component = self.getNextComponent()  
    self.in_service.append(next_component)  
    """ schedule departure for next component"""  
    depart = self.scheduleDeparture(next_component, False)  
  
    return component, depart
```

Here is a basic flowchart showing the execution loop for the Inspector class:



Workstation Class

The workstation class is used to represent the workstations within the simulation, and like the Inspector class, includes two primary methods: put, and get. Workstation deviates from Inspector in the additional logic necessary for managing the buffers, and determining whether or not the workstation has the necessary components to begin service on a product, thus generating a departure event. If the necessary components are not yet in the buffers, then no event is generated. Otherwise, workstation works similarly to Inspector in that it will attempt to add a new product to service when one is removed if the buffers contain those components required. Here is a flow chart showing the logic for a component arrival to a workstation:




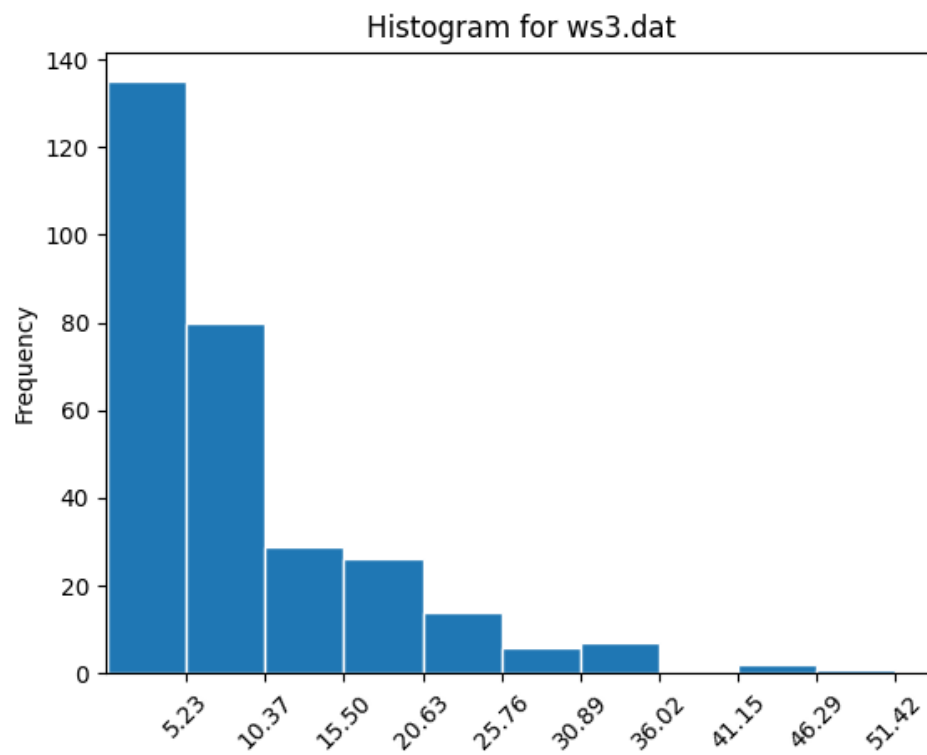

Model Translation

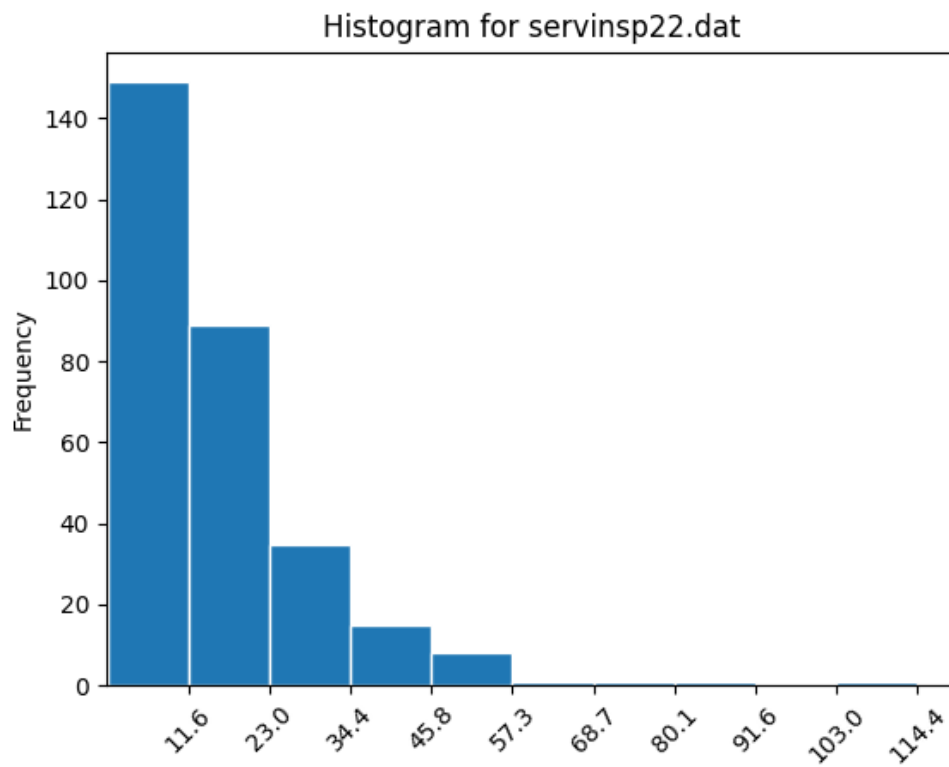
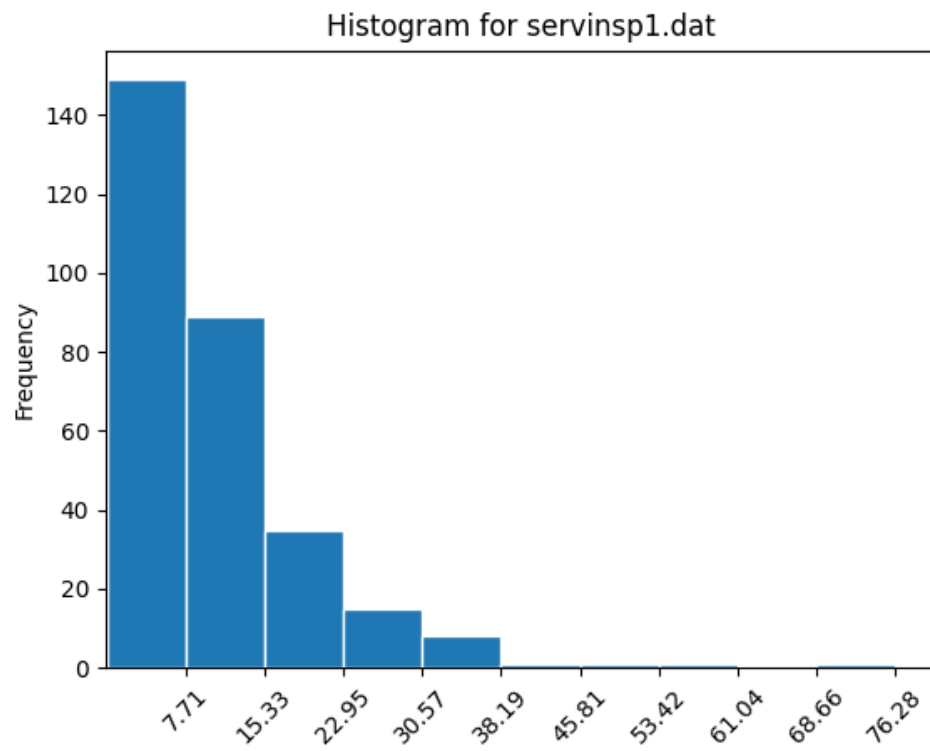
The implementation of the discussed model will be created using python. A general purpose, high level programming language was chosen due to the relative simplicity of the simulation, along with the low cost. Additionally, the team developing the simulation have experience with python and thus no time will be wasted learning a new tool. Python also allows for the simulation to be flexible, and alternative configurations to be easily tested and verified.

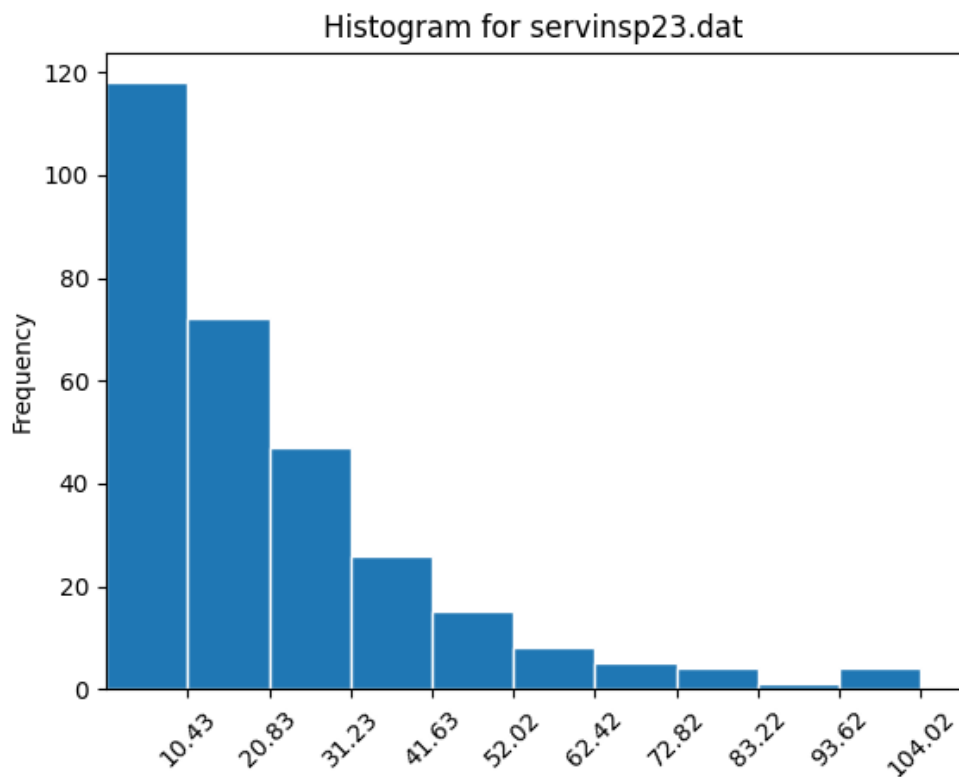
Data Collection Input Modelling

First we generated histograms for each file containing historical data for the 2 inspectors and 3 workstations. The bin size was chosen using Sturges's rule, which is a common method used to determine the appropriate number of bins for a histogram^[1]. Here are the resulting histograms:









These however are not fully adequate for input modelling, since a histogram just provides the general shape of the distribution of data. We need to employ more rigorous statistical techniques like QQ plots and chi-square tests to verify that they actually follow a theoretical statistical distribution that can be estimated from a histogram.

QQ plots

QQ plots are a tool used for comparing two distributions. In this section we are going to use these plots to compare our measured data against a distribution to determine if our data fits that distribution closely. Based on the nature of the data being the times between events (arrival and departure), and the shape of the histograms, we will assume an exponential distribution, and compare our values against it.

A QQ plot works by comparing the quantiles of our measured data versus the distribution. If the quantiles are close, then the plot should show a slope close to 1. To do this we need to first obtain the inverse cdf of the theoretical distribution, and calculate it for each point, where q is equal to the quantile of our measured data for that point. Then, by calculating the inverse cdf of the distribution with respect to the quantile, we should get a value close to the measured value.

To get the slope expected, we need to calculate each quantile such that it is with respect to our values in ascending order.

Knowing the cdf of an exponential distribution is:

$$1 - e^{-\lambda q}$$

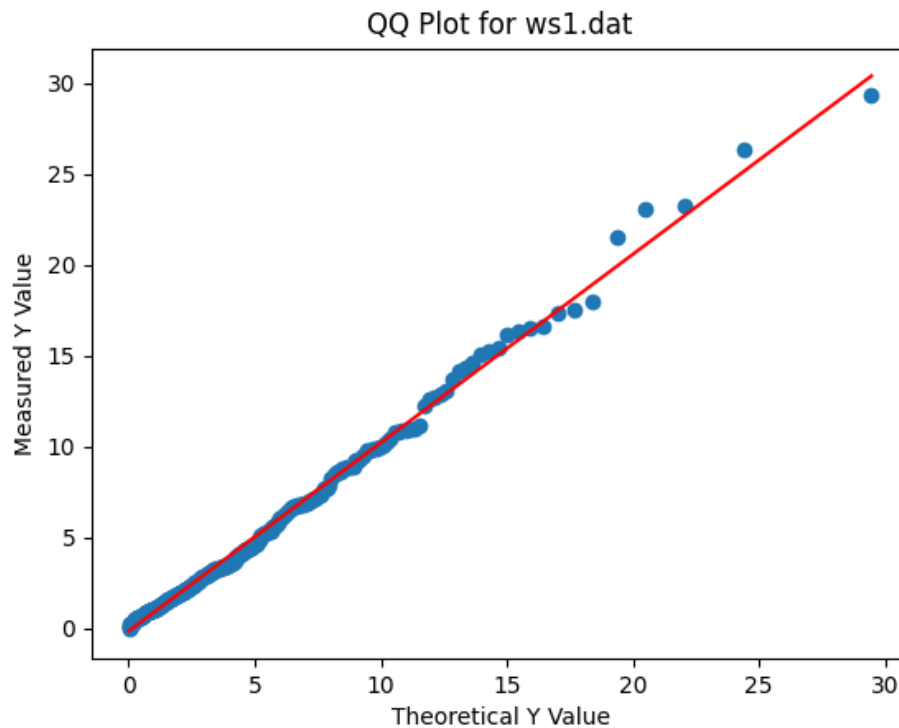
We can calculate the inverse cdf to be:

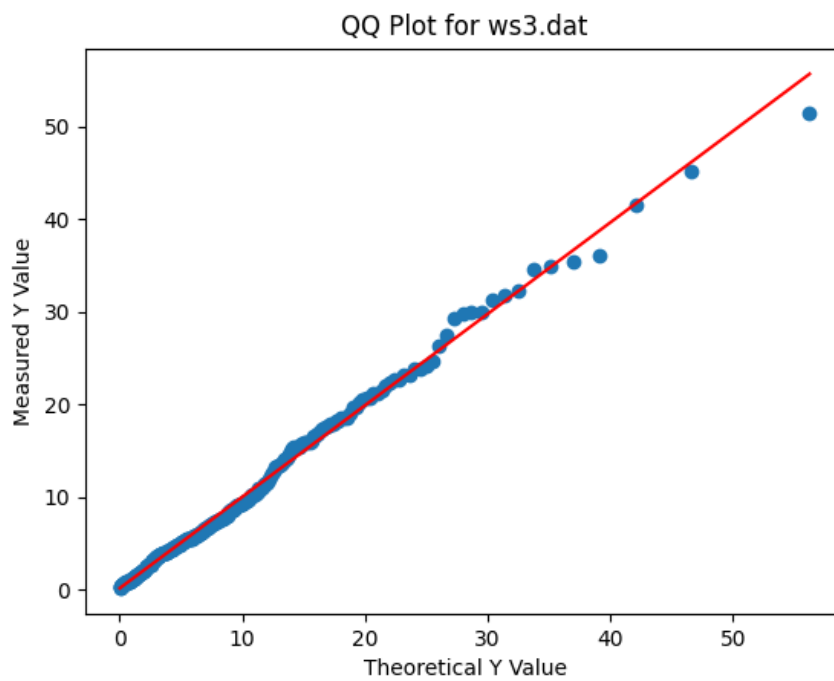

$$- \ln(1 - q)/\lambda$$


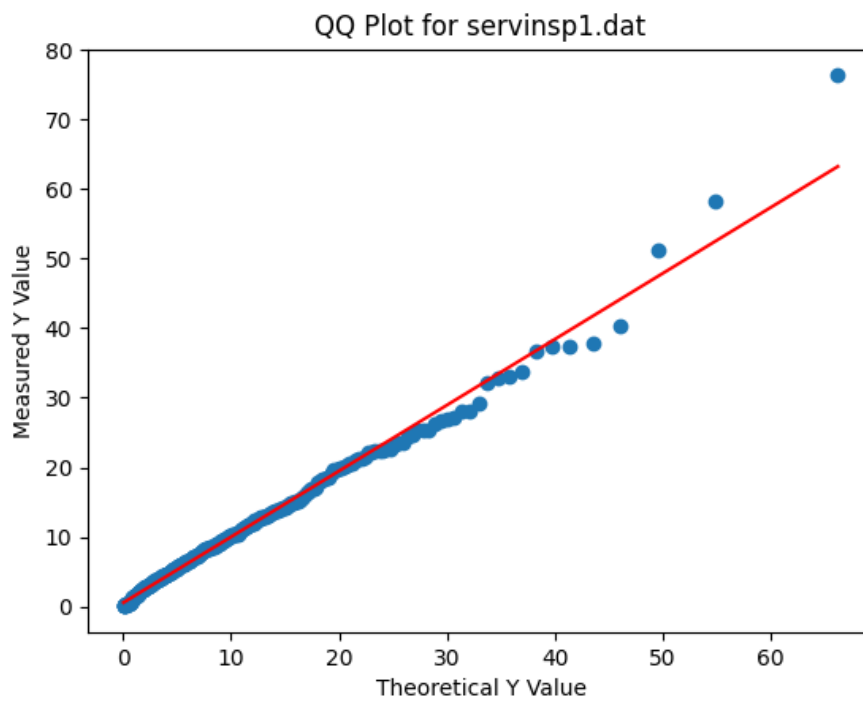
Where $q = j - \frac{0.5}{n}$, knowing that j is equal to the sequence number of the value in our ordered list of values, and n is the total number of values.


In order to calculate the values, we need to estimate what the rate parameter λ . To do this, we will use the maximum likelihood estimation, which leaves the recommended estimation for the rate parameter of an exponential distribution as $1 / X$, where X is the mean of the data set. This means that for each QQ plot, the rate parameter will be unique. The resultant plots are as follow where x is the theoretical quantiles and y is the sample quantiles:

Plots:









QQ Plot Conclusion

It can be seen that the plots all follow a fairly straight line. QQ plots offer a visual indicator of a distribution's fit, but there is no metric for precise evaluation. Additionally, it can be noted that the plots appear to deviate from the straight line when approaching the tail, the variance for the higher values are higher, and as a consequence, the slope representing the middle of the data set is more useful for identifying the quality of the fit. Given this information, it can be seen that an exponential distribution results in a slope close to 1, and will consequently be used moving forward.

Chi Square Tests

Now that we have a robust reasoning for our hypothesis that the data follows an exponential distribution, we need to use a goodness of fit test to get a quantitative determination on whether it is adequate. To do this, a chi-square test will be used. The chi square test is used to analyse the difference between our expected values and observed ones, specifically using the number of values that we expect to fall in some intervals for each distribution. This test relies on the central limit theorem, which states that the mean of independent random variables tends towards a normal distribution, even if the original random variables follow a different one, like in our case, potentially an exponential one. A random variable with the chi square distribution is a variable representing the sum of some number of these independent normal random variables.

Therefore, knowing that the following equation produces an approximately normal random variable:

$$\chi = \frac{(O_i - E_i)^2}{E_i}$$

Where O_i is our observed number of values, and E_i is the theoretical number of values for a given interval. We can sum these values for each interval, and check the result against the chi-squared distribution.

Sample Calculation for Servinsp1:

Hypothesis Formulation

Let our hypotheses be:

$$H_0: \text{Sample data follows exponential distribution}$$

$$H_1: \text{Sample data does not follow exponential distribution}$$

Determination of k value

The k value for a data set represents the number of intervals to be used and is determined based on the number of samples. We have $n = 300$ samples for each data set, and it is recommended that k be between the root of n and $n / 5$ for $n > 100$. For this section we have selected $k = 20$, which falls in the recommended range.

Determination of Probability

Each interval should have an equal probability. Given $k = 20$:

$$p = 1/k$$
$$p = 1/20 = 0.05$$

Determination of class intervals

For a given k, and knowing the mean of servinsp. = μ , the interval is determined by:

$$\alpha_i = -\frac{1}{\lambda} \ln(1 - kp)$$
$$\alpha_i = -\mu \ln(1 - kp)$$

For the first $k = 0$:

$$\alpha_0 = -\mu \ln(1)$$
$$\alpha_0 = 0$$

For the second $k = 1$:

$$\alpha_1 = -\mu \ln(1 - 1(0.05))$$
$$\alpha_1 = 0.5313$$

Therefore, the first interval will be from 0 to 0.5313. This then needs to be repeated for k intervals.

Determination of observed value for interval 1

The observed value for the first interval is the number of values from the data set that fall within that interval. For servinsp1.dat, 16 values fall between 0 and 0.5313, meaning $E_1 = 16$.

Determination of expected value for interval 1

The intervals represent equal probabilities, and as such we expect all the data values to be evenly divided among them. Therefore, the expected value for interval 1 is the same as the expected value for all other intervals, and is n / k , or $300 / 20 = 15$, meaning $O_1 = 15$.

Determination of Test Statistic

$$\chi_0^2 = \sum_{i=1}^k \frac{(O_i - E_i)^2}{E_i}$$

Sample Calculation for interval 1

$$\frac{(O_i - E_i)^2}{E_i} = \frac{(15 - 16)^2}{16} = 0.0625$$

This then needs to be done for each interval, and the results summed. The individual values for each calculated value, and their sums can be found in the appendix of this report. The sum for the results of the calculations for all intervals of servinsp1.dat has been determined to be 16.6667.

Comparison of calculated value versus expected statistic for servinsp1

$$\chi_0^2 = 16.667$$
$$\chi_{0.5,18}^2 = 28.869$$

Since $\chi_{0.5,18}^2 > \chi_0^2$ the null hypothesis is not rejected, and the exponential distribution is determined to be a good fit with the confidence interval of 0.05.

Chi-Squared tests for other data sets

The individual intervals and calculated values are provided for each data set in the appendix. Here we will simply list the resultant calculated statistic for each set.

$$\text{Servinsp22: } \chi_0^2 = 16.667$$

$$\text{Servinsp23: } \chi_0^2 = 16.933$$

$$\text{W1: } \chi_0^2 = 24.8$$

$$\text{W2: } \chi_0^2 = 20.8$$

$$\text{W3: } \chi_0^2 = 21.333$$

It can be seen that for all of these data sets, $\chi_{0.5,18}^2 > \chi_0^2$, which means that the null hypothesis will be accepted for each.

Input Generation

Now that we have completed our input modelling, we can generate random numbers that adequately mirror the actual behaviour of the system being modelled by using the historical data that was collected. It is impossible for computers to generate truly random numbers, so we must create a way in order to generate random numbers that are uniform and independent. This means that every random number has an equal chance of being selected, and each random number is not affected by future or previous values in the sequence. In addition to uniformity and independence, two important properties of a good random number generator are maximum density and maximum period. This means that our generator should not produce clusters of similar values, which is important because if the RNG is not densely populated, it can lead to biased results and affect the accuracy of the simulation or other application. The generator should also have a large period, which is the number of unique values that can be produced before there is repetition. This is important because if the period is too short, the RNG may start to produce the same values again, leading to repeating patterns.

Implementation

We chose to use a CLCG generator, because LCG has a period which isn't not adequate for complex system simulation like our project. The implementation of CLCG in code involves setting up multiple LCGs with different parameters, and then combining their outputs to produce the final sequence of random numbers. Each LCG is typically defined by the following formula:

$$X_{i+1} = (aX_i + c) \bmod m, i = 0, 1, 2, \dots$$

where X_i is the current value in the sequence, a is a multiplier, c is an increment, and m is the modulus. When then need to combine the outputs from each of our LCG's using the following formula:

$$X_i = \left(\sum_{j=1}^k (-1)^{j-1} X_{i,j} \right) \bmod m_1 - 1$$

To implement CLCG, we first choose several sets of values for the parameters a , c , and m , and then create a separate LCG for each set of parameters. We can then use the outputs of these LCGs to produce the final sequence of random numbers. This can be seen in our random function, which returns our random number after performing the preceding equations:

```
def random(self):  
    # Generate two random integers using the two LCGs  
    self.x1 = (self.a1 * self.x1) % self.m1  
    self.x2 = (self.a2 * self.x2) % self.m2  
    z = (self.x1 - self.x2) % (self.m1 - 1)
```

In terms of choices for our constants, L'Ecuyer suggests using two LCGs with prime moduli and with the values of the parameters chosen such that the LCGs have a long period and good statistical properties^[2]. In particular, L'Ecuyer recommends using the following values:

$a_1 = 40,014,$
 $m_1 = 2,147,483,563,$
 $a_2 = 40,692,$
 $m_2 = 2,147,483,399.$
 $c_1 = c_2 = 0$

Uniformity Testing


In order to confirm the uniformity of our random number generator, we must test our hypothesis of:

$H_0: X \text{ conforms to a uniform distribution}$

$H_1: X \text{ does not conform to a uniform distribution}$

To perform chi-square testing on our randomly generated numbers to determine how well they fit a uniform distribution, we must first divide our data into class intervals. Since our sample size

is 300, we can choose a number of class intervals in the range $n^{1/2}$ to $n/5$. The same as before, we will choose 20. Since the expected distribution is uniform, our expected frequency is $300/k$ which equals 15. The histogram representing this grouping is:



We can then calculate the chi-squared test statistic using this equation:

$$\chi_0^2 = \sum_{i=1}^k \frac{(O_i - E_i)^2}{E_i}$$

Doing this for our histogram, we get the chi-squared test statistic value of 18.53. We then need to compare this to $\chi_{0.5,18}^2 = 28.869$.

Since $\chi_{0.5,18}^2 > \chi_0^2$ the null hypothesis is not rejected, and the exponential distribution is determined to be a good fit for a uniform distribution with the confidence interval of 0.05.

Independence Testing

In order to verify for the independence of our random number generator, we must test our hypothesis of:

$$H_0: \rho_{il} = 0$$

$$H_1: \rho_{il} \neq 0$$

In order to calculate this, we must first determine $\hat{\rho}_{il}$ by using:

$$\hat{\rho}_{il} = \frac{1}{M+1} \left[\sum_{k=0}^M R_{i+kl} R_{i+(k+1)l} \right] - 0.25$$

We must also find $\hat{\sigma}_{\rho il}$ using:

$$\hat{\sigma}_{\rho il} = \frac{\sqrt{13M+7}}{12(M+1)}$$


Doing this, we can calculate our test statistic by dividing these two quantities. From Table A.3, we can determine our z score for 0.025 confidence interval to be 1.96. If our z score from our autocorrelation test is less than this, we fail to reject the null hypothesis.

$\alpha = 0.05$, $i = 5$, $l = 5$ $N = 300$, and $M = 58$

$$\hat{\rho}_{5,5} = \frac{1}{58+1} \left[\sum_{k=0}^{58} R_{5+5k} R_{5+5(k+1)} \right] - 0.25 = -0.16494$$

$$\hat{\sigma}_{\rho il} = \frac{\sqrt{13*58+7}}{12(58+1)} = 0.03896$$

Diving these two together, we get $-0.16494 / 0.04316 = -4.2336$, therefore we fail to reject the null hypothesis. I also had this test repeated for every lag value from 1 to 25 in order to find any evidence of correlation. As we can see, there isn't any significant deviation through the lag values, so we can conclude that the values are not correlated.



RVG Generation


In order to generate the appropriate RVG using the CDF technique, we must first use our random number generator to produce a value between 0 and 1. This value must be from a uniform distribution, which we confirmed in the previous step. Since our target distribution is exponential, we can calculate the inverse cumulative distribution function using:

$$F^{-1}(U) = -\frac{\log(1-U)}{\lambda}$$

Where lambda represents the rate parameter of the exponential distribution. This was estimated for each sample file earlier, so it can be used in our RVG for the main simulation. Implementing this in code is fairly simple where U is our uniform random number and y is lambda:

```
return -np.log(1 - U) / y
```

Running this function on our random uniform distribution we generated earlier, we can create the following histogram using a lambda value of 1 as an example which follows an approximately logarithmic distribution:



Verifying a model is an important step in the simulation process. It ensures that the model accurately reflects the real-world system being studied and that it is reliable for making predictions and drawing conclusions. There are several validation techniques that can be used to assess the accuracy and reliability of the simulation model.

One method that can be used to validate the system is to check the output for reasonableness. This involves looking at specific variables to ensure that they fall within a reasonable range. For example, our average buffer occupancies for each workstation's buffers must be between 0 and 2. Similarly, the throughput cannot be negative, and the average components in the system cannot be less than 2 and more than 15. The values can be seen throughout the report as they are all used throughout the statistical analysis.

Furthermore, another validation technique that can be used is to perform a trace of the simulation run. This involves examining the current contents and total counts of variables throughout the simulation run. Analytical results can also be used to validate the model. This includes comparing the simulation results using statistical tests like Little's Law, and confidence interval testing which will be seen later in this report.

These validation techniques can be used to assess the accuracy and reliability of a simulation model. By using a combination of these techniques, it is possible to ensure that the model is robust and can be used to make informed decisions.

Code Verification by Little's Law

Little's Law relates the number of entities in the system (L) to the arrival rate and the amount of time each entity spends in the system. To verify the system is behaving as expected, we have evaluated little's law for two levels of granularity. First for the entire system, including the inspectors, workstations, and their buffers, and then for each individual workstation and its associated buffers.

If the system is behaving as expected, the average number of components in each of these systems should equal to the arrival rate multiplied by the amount of time each component spends in the system, barring some slight variation caused by rounding error. To get the values, the system was run for eight replications, and the values were only measured beyond the deletion point as determined by initialization analysis in the following sections. The arrival rate was calculated by measuring the total number of components that arrive at each inspector, or buffer divided by the time. However, it could also have been measured by analyzing the throughput of the components that left the system, as, for a steady state system throughput is equal to the arrival rate.

Entire Facility

- Arrival Rate: 0.10765805429501399
- Amount of time component Spends in System: 51.6897127104437
- Average number of entities in system: 5.572567720300621

Little's Law: $0.107658 \times 51.6897 = 5.572568$

$5.5648 = 5.5727$

Workstation 1

- Arrival Rate 0.08525171777719182
- Amount of time component Spends in System 6.044015773077008
- Average number of entities in system 0.5152603555791313

Little's Law: $0.08525 \times 6.0440 = 0.51526$

$0.51526 = 0.51526$

Workstation 2

- Arrival Rate 0.011252449842275126
- Amount of time component Spends in System 96.44537860522436
- Average number of entities in system 1.0836005232596713

Little's Law: $0.011252 \times 96.445379 = 1.0836005232596713$

$1.08543 = 1.0836$

Workstation 3

- Arrival Rate 0.011134323374336744
- Amount of time component Spends in System 177.7061792973173
- Average number of entities in system 1.973706841461812

Little's Law: $0.0111343 \times 177.706 = 1.97371$

$1.9786 = 1.9737$

It can therefore be seen that the values calculated through the use of little's law match very closely to the values obtained through measuring the results of the simulation.

In order to show that the number of replications is enough, we can use an operating characteristic curve to find the minimum sample size needed to detect that effect size with a certain level of statistical power and significance level. Statistical power refers to the ability of a test to detect a true effect when it is present, while the significance level is the probability of rejecting the null hypothesis when it is true. By choosing an appropriate number of replications, we can increase the power of our analysis and reduce the likelihood of making a type II error (i.e., failing to detect a true effect).


Estimating $\hat{\sigma} = S = 0.0556209$. This was determined from our simulation.

Specify the critical difference $|E(Y_2) - \mu_0| = 0.082$

Calculate:

$$\hat{\delta} = \frac{0.082}{0.0556209} = 1.47427$$

Find our sample size n by look up operating characteristic curve with $\alpha = 0.05$ and $\beta = 0.1$. The OC curve plots the probability of correctly rejecting the null hypothesis (i.e., detecting a true difference) against the sample size for different effect sizes:



Therefore we can conclude that $n \approx 7$ replications is a good starting point for our analysis.

We assume that there is a known output performance measure for the existing system. This was provided in the following table as the average buffer occupancy for each workstation:

C1 Buffer of Workstation 1	C1 Buffer of Workstation 2	C2 Buffer of Workstation 2	C1 Buffer of Workstation 3	C3 Buffer of Workstation 3
0.28	0.41	0.60	0.32	1.75

This value will be denoted by μ_0 , and our unknown performance measure of the simulation will be denoted by μ . The confidence-interval formulation tries to bound the difference $|\mu - \mu_0|$ to see whether it is $\leq \varepsilon$. The requirement that confidence intervals do not exceed 20% of the estimated value can only be controlled by the number of replications, so we'll do $20\% * \text{the average buffer occupancy for each buffer}$. For C1 buffer of workstation 2 this is 0.41

$* 0.2$ which is 0.082. The results for each buffer from the simulation for 6 replications are as follows:

W1:C1 Average Buffer Occupancy: 0.12205

W1:C1 Standard Deviation: 0.00571

W2:C1 Average Buffer Occupancy: 0.43097

W2:C1 Standard Deviation: 0.05562

W2:C2 Average Buffer Occupancy: 0.56418

W2:C2 Standard Deviation: 0.08180

W3:C1 Average Buffer Occupancy: 0.01823

W3:C1 Standard Deviation: 0.00695

W3:C3 Average Buffer Occupancy: 1.86116

W3:C3 Standard Deviation: 0.03838

Then, if Y is the simulation output and $\mu = E(Y)$, then we execute the simulation n number of replications and form a confidence interval for μ , such that $\bar{Y} \pm t_{\alpha/2, n-1} S/\sqrt{n}$. Here is the sample calculation for 6 replications on W2 C1 buffer:

$$\bar{Y} \pm t_{\alpha/2, n-1} S/\sqrt{n} \Rightarrow 0.4309748 \pm t_{0.05/2, 6-1} 0.0556209/\sqrt{6}$$

Using a t-distribution table or a calculator, we can find that $t_{0.025, 6-1} = 2.447$ for a double sided t-test.

$$0.4309748 \pm 2.447 * 0.0556209/\sqrt{6} = (0.37541, 0.486539)$$

Therefore, with 95% confidence, the true population mean μ is expected to fall within the range of 0.37541 – 0.486539. Our confidence interval contains μ_0 , so we must find the best and worst case error next. The best case error is $|0.41 - 0.37541| = 0.0346$, and the worst case error is $|0.4865 - 0.41| = 0.0765$. Our worst case error of 0.0765 is less than our ε value of 0.082, so we can conclude that the simulation model is close enough to be considered valid at 6 replications.

This procedure was repeated for the other buffers. We got acceptable confidence range measurements for w2c2 and w3c3, but not for w1c1 and w3c1. We were told that only the procedure would be marked, not numbers because the data is sensitive and can be quite variable with minor implementation changes, we will stop at 6 replications. The confidence intervals for the other 4 buffers are as follows:

$$W1\ C1: 0.12205 \pm 2.447 * 0.0057075/\sqrt{6} = (0.11635, 0.12775)$$

$$W2\ C2: 0.56418 \pm 2.447 * 0.081799/\sqrt{6} = (0.48246, 0.645896)$$

$$W3\ C1: 0.018231 \pm 2.447 * 0.0069501/\sqrt{6} = (0.011288, 0.025174)$$

$$W3\ C3: 1.8612 \pm 2.447 * 0.038379/\sqrt{6} = (1.82286, 1.89954)$$

To determine the initialization phase, a Welch plot can be used. By calculating the average of a performance metric across all replications, and plotting it, one can see where the data starts to vary around a consistent mean. It is after this point that the data can start to be collected. For the Welch plot, the time for a component to pass through the system was considered as the performance metric, and measured against the total

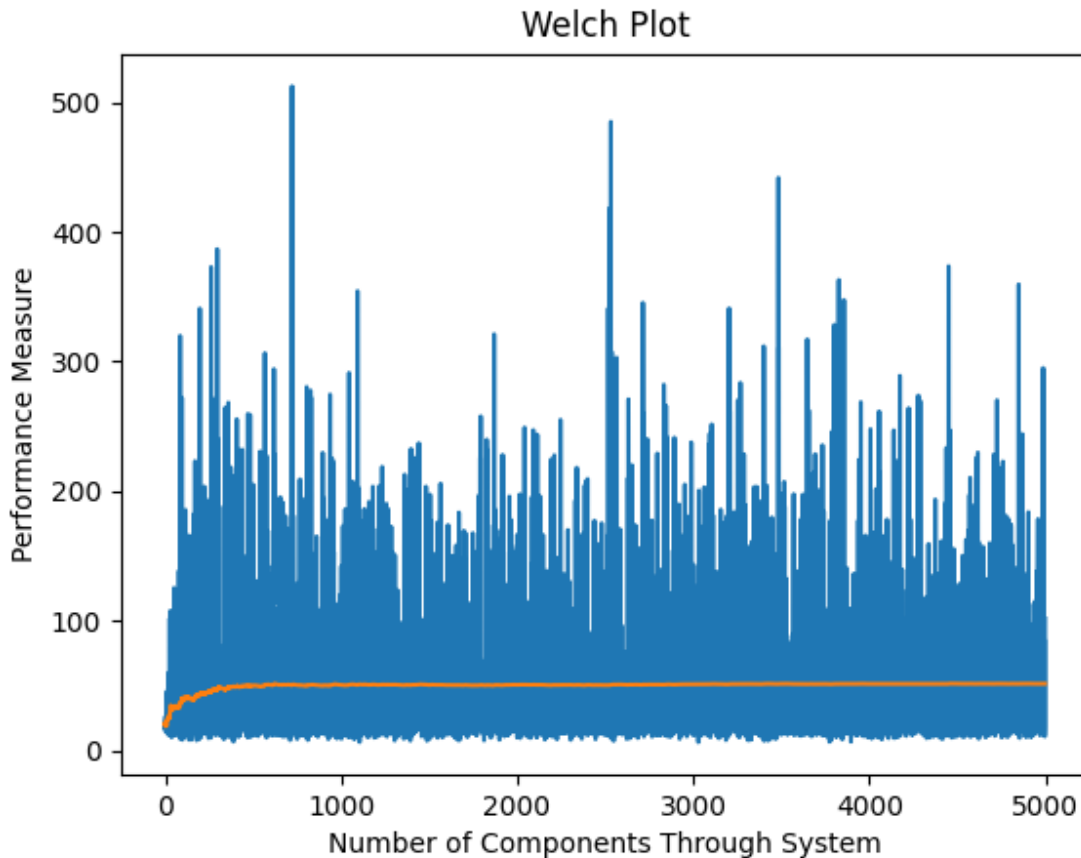
number of components that have passed through the system to represent the time that the system has been alive. Six replications were adequate to view the initialization period.

The averages plotted were determined using the following equation:

$$Y.j = \frac{1}{R} \sum_{r=1}^R Y_{rj}$$

Where $Y.j$ represents the average of the value across all replications.

The data can be noisy, so to smooth it, a cumulative average of the Welch averages can be calculated for each point and plotted alongside. This is represented by the orange line below, and shows the initialization period clearly.



From this data, we can determine that an initialization phase of about 1000 components having passed through the system is adequate, therefore we will set $d = 1000$. Knowing

that T_e should be at least 10 times greater than T_d , the phase for data collection will be set to be 10 000

Replication Method - Sample Calculation for Component Time in System

First step: get averages of all non-deleted data, for each replication.

$$Y_r(n, d) = \frac{1}{(n-d)} \sum_{j=d+1}^n Y_{rj}$$

For replication 1.

Replication	$Y_r(n, d)$
1	52.023904404742765
2	51.908259952716485
3	51.14631579340045
4	51.41186974742246
5	52.817739924031294
6	51.606831811349075

Calculate overall point estimator

$$Y_{..}(n, d) = \frac{1}{R} \sum_{r=1}^R Y_r(n, d)$$

$$= 51.81915361$$

Calculate standard variance

$$S^2 = \frac{1}{R-1} \left(\sum_{r=1}^R (Y_r - Y_{..})^2 \right)$$

$$S^2 = \frac{1}{6} \left(\sum_{r=1}^6 (Y_r - 51.81915361)^2 \right)$$

$$= 0.2851181455$$

Calculate Standard Error

$$s.e(Y..) = \frac{s}{\sqrt{R}} = 0.5339645546 / \sqrt{6} = 0.2179901166$$

Using 95% confidence, alpha = 0.5, and R - 1 = 5

$$t_{0.025,5} = 2.571 \text{ leaving the confidence interval}$$

$$51.2595 < w < 52.3804$$

Replications required to reduce the interval to range of 0.2

$$R \geq \left(\frac{z_{0.025} S_0}{0.2} \right)^2 = \left(\frac{1.96 * 0.2851181455}{0.2} \right)^2 = 7.8$$

Therefore, 8 replications would be required

Alternative Operating Policy

The current system spends much of its time blocked, as W1 uses so many C1 components, W2 and W3 often have long wait-times to receive the components they need from Inspector 1. In order to solve that, we propose a policy wherein, instead of providing C1 components to the buffer with the least current components, the system simply provides components to the buffer's in order of w1, w2, and w3, before rolling back to w1. This ensures that all of the workstations get an equal number of C1 buffers, and no buffer is starved.

To compare the difference in the system for each alternative, let's consider the percent of time each of the inspectors spent blocked. The old operating time resulted in the following:

```
Inspector 1 blocked times 0.0
Inspector 2 blocked times 0.9088432956904341
```

Now, let's consider the new operating policy with respect to blocked times.

```
Inspector 1 blocked times 0.0
Inspector 2 blocked times 0.6395168591336154
```

To compare these results, we can take the confidence interval of the difference.

$$\bar{Y}_{.1} - \bar{Y}_{.2} \pm t_{\alpha/2, \nu} s.e.(\bar{Y}_{.1} - \bar{Y}_{.2})$$

In order to do this, we must calculate the standard error, and the degrees of freedom for the data. Where standard error is given by:

$$s.e.(\bar{Y}_{.1} - \bar{Y}_{.2}) = \sqrt{\frac{S_1^2}{R_1} + \frac{S_2^2}{R_2}}$$

And degrees of freedom are given by:

$$\nu = \frac{(S_1^2 / R_1 + S_2^2 / R_2)^2}{\left[(S_1^2 / R_1)^2 / (R_1 - 1) \right] + \left[(S_2^2 / R_2)^2 / (R_2 - 1) \right]},$$

For data collection, 8 replications were used, and the means of the blocked time for inspector two are 0.90844 and 0.6395 respectively. Using these values, this results in the following standard error and degrees of freedom:

This results in a standard error of 0.39278, and 81 degrees of freedom. That leaves the t value to be: 1.990 for an alpha of 0.05.

These leaves the confidence interval to be

$$0.268941 \pm 0.10563487$$

Lower bound difference results in 0.163, which equates to a 17.97% reduction in blocked time for inspector 2.

Appendix

Chi-Square Calculations

Servinsp1

interval	actual	expected	calculation
0.0 - 0.5312913268 697537	16	15	0.0666666666 7
0.5312913268 697537 - 1.0913147387 373547	7	15	4.266666667
1.0913147387 373547 - 1.6833564450	11	15	1.066666667

Chase Scott - Emmitt Luning

342968			
1.6833564450 342968 - 2.3113008215 929645	10	15	1.666666667
2.3113008215 929645 - 2.9797850150 69024	15	15	0
2.9797850150 69024 - 3.6944069685 724337	14	15	0.066666666 7
3.6944069685 724337 - 4.4620106744 23191	18	15	0.6
4.4620106744 23191 - 5.2910858366 61989	16	15	0.066666666 7
5.2910858366 61989 - 6.1923418484 96644	19	15	1.066666667
6.1923418484 96644 - 7.1795561129 93658	17	15	0.266666667
7.1795561129 93658 - 8.2708708517 31015	15	15	0
8.2708708517 31015 - 9.4908569345 86626	19	15	1.066666667
9.4908569345 86626 - 10.873963081	18	15	0.6

Chase Scott - Emmitt Luning

566092			
10.873963081			
566092 -			
12.470641949			0.0666666666
655649	14	15	7
12.470641949			
655649 -			
14.359112225			
987316	21	15	2.4
14.359112225			
987316 -			
16.670413047			
580283	13	15	0.2666666667
16.670413047			
580283 -			
19.650198062			
649313	13	15	0.2666666667
19.650198062			
649313 -			
23.849969160			
57394	20	15	1.666666667
23.849969160			
57394 -			
31.029525273			
56761	12	15	0.6
>			
31.029525273			
56761	12	15	0.6
	300	300	16.66666667

Servinsp22

interval	actual	expected	calculation
0.0 -			0.0666666666
0.7969389565	16	15	7

Chase Scott - Emmitt Luning

47582			
0.7969389565 47582 - 1.6369761469 257988	7	15	4.266666667
1.6369761469 257988 - 2.5250408974 437426	11	15	1.066666667
2.5250408974 437426 - 3.4669597862 255803	10	15	1.666666667
3.4669597862 255803 - 4.4696885504 16313	15	15	0
4.4696885504 16313 - 5.5416241253 98168	14	15	0.0666666666 7
5.5416241253 98168 - 6.6930325249 79904	18	15	0.6
6.6930325249 79904 - 7.9366483366 41895	16	15	0.0666666666 7
7.9366483366 41895 - 9.2885356898 29995	19	15	1.066666667
9.2885356898 29995 - 10.769360740 132408	17	15	0.266666667
10.769360740 132408 - 12.406336887	15	15	0

Chase Scott - Emmitt Luning

05821			
12.406336887 05821 - 14.236320526 357993	19	15	1.066666667
14.236320526 357993 - 16.310984865 53058	18	15	0.6
16.310984865 53058 - 18.706009076 774304	14	15	0.066666666 7
18.706009076 774304 - 21.538721480 264815	21	15	2.4
21.538721480 264815 - 25.005681266 4904	13	15	0.266666667
25.005681266 4904 - 29.475369816 90672	13	15	0.266666667
29.475369816 90672 - 35.775042006 62281	20	15	1.666666667
35.775042006 62281 - 46.544402746 75523	12	15	0.6
> 46.544402746 75523	12	15	0.6
	300	300	16.66666667

Servinsp23

Chase Scott - Emmitt Luning

interval	actual	expected	calculation
0.0 - 1.0583220617 3003	15	15	0
1.0583220617 3003 - 2.1738778818 42453	12	15	0.6
2.1738778818 42453 - 3.3532135260 547458	15	15	0
3.3532135260 547458 - 4.6040665960 01935	15	15	0
4.6040665960 01935 - 5.9356741982 599655	15	15	0
5.9356741982 599655 - 7.3591873273 84841	9	15	2.4
7.3591873273 84841 - 8.8882390838 92696	23	15	4.266666667
8.8882390838 92696 - 10.539740794 261903	14	15	0.0666666666 7
10.539740794 261903 - 12.335025362 92053	9	15	2.4
12.335025362 92053 -	14	15	0.0666666666 7

Chase Scott - Emmitt Luning

14.301537110 679414			
14.301537110 679414 - 16.475414992 52187	18	15	0.6
16.475414992 52187 - 18.905603706 681358	18	15	0.6
18.905603706 681358 - 21.660724438 064257	21	15	2.4
21.660724438 064257 - 24.841277904 941318	14	15	0.0666666666 7
24.841277904 941318 - 28.603074221 358828	10	15	1.666666667
28.603074221 358828 - 33.207140817 36077	20	15	1.666666667
33.207140817 36077 - 39.142815015 62075	15	15	0
39.142815015 62075 - 47.508677928 040186	14	15	0.0666666666 7
47.508677928 040186 - 61.810215038 71962	14	15	0.0666666666 7
> 61.810215038	15	15	0

Chase Scott - Emmitt Luning

71962			
	300	300	16.93333333

W1

interval	actual	expected	calculation
0.0 - 0.2361756995 6627757	6	15	5.4
0.2361756995 6627757 - 0.4851237143 034894	12	15	0.6
0.4851237143 034894 - 0.7483048676 283796	18	15	0.6
0.7483048676 283796 - 1.0274458867 303353	21	15	2.4
1.0274458867 303353 - 1.3246081290 981868	18	15	0.6
1.3246081290 981868 - 1.6422800564 53898	17	15	0.2666666667
1.6422800564 53898 - 1.9835040385 713643	19	15	1.066666667
1.9835040385 713643 - 2.3520540158 28523	16	15	0.0666666666 7

Chase Scott - Emmitt Luhning

2.3520540158 28523 - 2.7526906502 2919	13	15	0.2666666667
2.7526906502 2919 - 3.1915384306 2322	15	15	0
3.1915384306 2322 - 3.6766621449 2671	23	15	4.2666666667
3.6766621449 2671 - 4.2189843173 53557	11	15	1.0666666667
4.2189843173 53557 - 4.8338184870 77118	13	15	0.2666666667
4.8338184870 77118 - 5.5435924464 51743	11	15	1.0666666667
5.5435924464 51743 - 6.3830768612 4644	10	15	1.6666666667
6.3830768612 4644 - 7.4105227479 76776	20	15	1.6666666667
7.4105227479 76776 - 8.7351308770 74965	9	15	2.4
8.7351308770 74965 - 10.602061178 599996	17	15	0.2666666667

Chase Scott - Emmitt Luning

10.602061178 599996 - 13.793599609 223222	13	15	0.2666666667
> 13.793599609 223222	18	15	0.6
	300	300	24.8

W2

interval	actual	expected	calculation
0.0 - 0.5689763392 786393	19	15	1.0666666667
0.5689763392 786393 - 1.1687227583 894413	20	15	1.6666666667
1.1687227583 894413 - 1.8027585608 06548	17	15	0.2666666667
1.8027585608 06548 - 2.4752436449 316777	17	15	0.2666666667
2.4752436449 316777 - 3.1911440747 591073	8	15	3.2666666667
3.1911440747 591073 - 3.9564548609 677423	10	15	1.6666666667
3.9564548609 677423 - 4.7785054469	18	15	0.6

Chase Scott - Emmitt Luning

33266			
4.7785054469 33266 - 5.6663877196 90786	17	15	0.2666666667
5.6663877196 90786 - 6.6315707001 61799	15	15	0
6.6315707001 61799 - 7.6888090360 604515	13	15	0.2666666667
7.6888090360 604515 - 8.8575317944 49894	19	15	1.0666666667
8.8575317944 49894 - 10.164052680 992134	16	15	0.0666666666 7
10.164052680 992134 - 11.645263897 028196	15	15	0
11.645263897 028196 - 13.355196755 751239	14	15	0.0666666666 7
13.355196755 751239 - 15.377618072 120903	17	15	0.2666666667
15.377618072 120903 - 17.852861717 052583	9	15	2.4
17.852861717 052583 - 21.044005791	8	15	3.2666666667

Chase Scott - Emmitt Luning

811696			
21.044005791 811696 - 25.541670753 113035	11	15	1.066666667
25.541670753 113035 - 33.230479789 1735	15	15	0
> 33.230479789 1735	22	15	3.266666667
	300	300	20.8

W3

interval	actual	expected	calculation
0.0 - 0.4511542742 4925236	7	15	4.266666667
0.4511542742 4925236 - 0.9267068443 096642	22	15	3.266666667
0.9267068443 096642 - 1.4294482459 120401	13	15	0.266666667
1.4294482459 120401 - 1.9626769570 682376	16	15	0.066666666 7
1.9626769570 682376 - 2.5303306828 154364	11	15	1.066666667
2.5303306828 154364 - 3.1371630034	11	15	1.066666667

Chase Scott - Emmitt Luhning

08638			
3.1371630034 08638 - 3.7889856011 24472	16	15	0.066666666667
3.7889856011 24472 - 4.4930076398 83675	21	15	2.4
4.4930076398 83675 - 5.2583231671 06123	18	15	0.6
5.2583231671 06123 - 6.0966314783 89447	22	15	3.2666666667
6.0966314783 89447 - 7.0233383226 99113	15	15	0
7.0233383226 99113 - 8.0593084354 57687	18	15	0.6
8.0593084354 57687 - 9.2337944817 98085	13	15	0.2666666667
9.2337944817 98085 - 10.589639118 273123	13	15	0.2666666667
10.589639118 273123 - 12.193262956 778893	9	15	2.4
12.193262956 778893 - 14.155939913	11	15	1.0666666667

847133			
14.155939913 847133 - 16.686270596 662574	16	15	0.0666666666 7
16.686270596 662574 - 20.252571392 236582	16	15	0.0666666666 7
20.252571392 236582 - 26.349202870 626037	17	15	0.2666666667
> 26.349202870 626037	15	15	0
	300	300	21.33333333

$X_{2,0.05} = 28.869$

References

- [1] Sturges, H. A. (1926). The choice of a class interval. Journal of the American Statistical Association, 21(153), 65-66.
- [2] P. L'Ecuyer, "Efficient and portable combined random number generators," Communications of the ACM, vol. 31, no. 6, pp. 742-749, June 1988
- [3] "Discrete time simulation inspector handling" YouTube, uploaded by Tadhg McDonald-Jensen, Feb 21, 2023, 40:48,
https://www.youtube.com/watch?v=KaG1uRvW2_A.